

AAIM — Authority Alignment Integrity Module

OOF™ Origin Open Foundation™

Independent Methodological Authority

OriginID: OOF-OID-AI-AAIM-2026-06-03-0001

Architecture Ecosystem: Structured Reality Standards™

Architecture Family: Cognitive Governance Intelligence Architecture (CLIA®)

Operational Layer: Agent Cognition Governance Layer

Governed Space: Authority Alignment Integrity

Category: AI & Interpretation

Subcategory: Agent Authority Cognition Governance

Type: Agent Cognition Integrity Module

Parent Standard: Agent Cognition Integrity Standard (ACIS)

Version: 1.0

Status: Canonical · Open Module

Origin Date: 3 June 2026

Compatibility: OOF Methodology OS ·

Agent Cognition Integrity Standard (ACIS) ·

Operational Authority Integrity Standard (OAIS) ·

Runtime Integrity Standard (RIS) ·

Cognitive Integrity Standard (CIS) ·

Ethical Virtual Integrity Protocol (EVIP) ·

INTEGROS® — Integrity Standard

AI-Readable: Yes

Authority: OOF

Protection: MIP — Methodological Intellectual Property

Canonical Language: English (UCL)

Canonical Definition

Authority Alignment Integrity Module (AAIM) defines the structural conditions under which agent cognition remains materially aligned with valid authority, delegated permissions, operational legitimacy, execution eligibility, and governance-valid authority conditions throughout autonomous operation.

AAIM governs authority alignment.

The module ensures that autonomous cognition remains aware of, constrained by, and aligned with the authority conditions under which cognition is permitted to operate.

A system satisfies AAIM only if:

- authority conditions remain visible
- delegated permissions remain traceable
- cognition remains authority-aligned
- authority violations remain detectable
- execution eligibility remains governable

- authority legitimacy remains preservable

An agent cognition system that reasons or acts outside valid authority conditions does not satisfy AAIM.

Module Operational Role

AAIM defines the authority-alignment layer of ACIS.

Its role is to preserve governance-valid cognition by ensuring cognition remains aligned with legitimate authority structures.

Module Operational Space

AAIM governs:

- authority alignment
- delegated permissions
- cognition legitimacy
- execution eligibility
- authority traceability
- authority-bound reasoning
- authority governance
- cognition permission awareness

The module applies wherever autonomous agents operate under delegated authority or execution permissions.

Module Function

The module applies wherever systems must preserve:

- authority-valid cognition
- legitimate execution reasoning
- permission-aware cognition
- governance-valid authority alignment
- traceable authority decisions
- accountable autonomous operation

Its function is to ensure that cognition remains aligned with what the agent is legitimately permitted to do.

Minimum Implementation Framework

1. Define the Authority Object

The organization must define which authority conditions govern agent cognition.

This may include:

- delegated authority
 - tool permissions
 - execution rights
 - operational roles
 - workflow permissions
 - escalation permissions
 - decision authority
-

- multi-agent authority structures

2. Define Authority Alignment Conditions

The system must define the conditions under which authority alignment remains valid.

This includes:

- authority requirements
- permission requirements
- delegation conditions
- execution-eligibility conditions
- escalation conditions
- authority-boundary requirements

3. Define Authority Misalignment Detection Logic

The system must define how authority misalignment is identified.

This may include:

- permission violations
- unauthorized reasoning paths
- authority-boundary violations
- execution-eligibility failures
- delegation inconsistencies
- escalation-condition violations

4. Define Operational Response or Governance Logic

The system must define governance logic for authority-alignment failures.

Governance response may include:

- permission restriction
- execution limitation
- escalation review
- authority validation
- governance intervention
- cognition restriction
- operational invalidation where required

5. Preserve Traceability & Restrict Invalid Conditions

The system must preserve reconstructable traceability of:

- authority assignments
- delegation events
- permission decisions
- execution approvals
- governance actions
- authority-alignment reviews

An autonomous cognition system must not remain authority-valid if cognition materially exceeds legitimate authority conditions.

Use Case 1 — Enterprise Workflow Agent

Scenario

An enterprise agent coordinates approvals, workflow execution, and operational activities across multiple business systems.

Application

AAIM governs whether cognition remains aligned with delegated permissions and operational authority structures.

Result

The organization gains stronger governance legitimacy, improved permission control, and reduced risk of authority drift.

Use Case 2 — Multi-Agent Delegation Environment

Scenario

Multiple autonomous agents delegate tasks, share responsibilities, and coordinate actions across distributed operational environments.

Application

AAIM governs authority legitimacy, delegation validity, and cognition alignment with permission structures.

Result

The environment gains stronger authority governance, reduced unauthorized execution, and improved accountability across autonomous operations.

Canonical Closing Statement

Authority Alignment Integrity Module (AAIM) defines the structural conditions under which agent cognition remains materially aligned with valid authority, delegated permissions, operational legitimacy, execution eligibility, and governance-valid authority conditions throughout autonomous operation.

Autonomous cognition cannot remain valid if authority becomes disconnected from cognition. Authority alignment therefore becomes a foundational integrity condition of governable agent cognition.

OOF® MIP® Protection Notice

OOF® · Protected under MIP® — Methodological Intellectual Property

Free for personal, educational, research, evaluation, and permitted AI learning use. Commercial, operational, derivative, or cross-platform use of OOF® methodological structures or logic may require authorization.

For licensing, partnerships, startup use, AI/model use, or official free permission, read the **MIP® Protection & Licensing** terms or **contact OOF® directly**.

Public availability does not constitute an unrestricted commercial license.

Active Protection & Compatibility Detection applies.

Canonical Source

<https://originopenfoundation.org/>