

Agent Cognition Integrity Standard (ACIS)

OOF™ Origin Open Foundation™

Independent Methodological Authority

OriginID: OOF-OID-AI-ACIS-2026-06-03-0001

Category: AI & Interpretation

Subcategory: Agent Cognition Governance Architecture

Type: Parent Standard

Version: 1.0

Status: Canonical · Open Standard

Origin Date: 3 June 2026

Compatibility: OOF Methodology OS ·

Cognitive Integrity Standard (CIS) ·

Human Cognition Integrity Standard (HCIS) ·

Runtime Integrity Standard (RIS) ·

Operational Authority Integrity Standard (OAIS) ·

Multi-Layer Truth Validation Framework (MTVF) ·

Ethical Virtual Integrity Protocol (EVIP) ·

INTEGROS® — Integrity Standard ·

Universal Canonical Language (UCL)

AI-Readable: Yes

Authority: OOF

Protection: MIP — Methodological Intellectual Property

Canonical Language: English (UCL)

Canonical Definition System

Canonical Definition

Agent Cognition Integrity Standard (ACIS) defines the structural conditions under which autonomous agent cognition remains valid, governable, traceable, bounded, interpretable, and operationally reliable during reasoning, planning, tool use, delegation, execution, adaptation, and multi-agent interaction.

ACIS governs agent cognition.

The standard ensures that autonomous agents remain cognitively valid as they increasingly perform tasks, make decisions, coordinate actions, access tools, interact with humans, and participate in operational environments.

Agent cognition increasingly depends not only on capability, autonomy, or execution performance, but on whether cognition remains governable throughout operation.

A. Standard Abstract

Future operational systems increasingly operate through:

- autonomous agents
- tool-using agents
- reasoning agents
- planning agents
- execution agents
- delegated agents
- multi-agent systems
- AI coworkers
- autonomous orchestration systems
- future agent ecosystems

Most existing governance approaches focus on:

- task completion
- agent performance
- execution success
- operational efficiency
- capability scaling

This creates a major governance problem.

An agent may:

- complete tasks successfully
- appear operationally effective
- produce useful outputs
- execute instructions correctly

while cognition underneath becomes:

- opaque
- unstable
- authority-confused
- tool-dependent
- execution-biased
- governance-invalid

An agent may therefore remain operationally useful while cognition itself becomes progressively less governable.

ACIS exists to govern that condition.

B. Core Principle

Agent cognition is not valid merely because an agent successfully executes tasks. Agent cognition becomes valid only when cognition remains traceable, governable, bounded, interpretable, authority-aligned, and operationally reliable throughout operation.

C. Scope

This standard may apply to:

- autonomous AI agents
- tool-using agents
- enterprise agents
- agentic AI systems
- multi-agent environments
- orchestration systems

- AI coworkers
- autonomous decision systems
- delegated execution systems
- future autonomous agent ecosystems

ACIS applies wherever autonomous agents participate in cognition-driven operational environments.

D. Why This Standard Exists

Autonomous agents increasingly perform activities previously reserved for humans.

Agents now:

- reason
- plan
- interpret
- select tools
- execute workflows
- coordinate with agents
- interact with humans
- make operational decisions

Traditional governance asks:

Can the agent perform the task?

ACIS asks a deeper cognition-governance question:

Does the agent remain cognitively governable while performing the task?

As autonomy increases, cognition itself becomes a governable operational space.

An agent may continue functioning while cognition underneath becomes progressively more difficult to audit, interpret, constrain, or supervise.

ACIS exists because agent cognition has become a distinct governance domain.

E. Agent Cognition Integrity Logic

Agent cognition integrity exists only when the following remain materially preservable:

1. Cognitive Traceability Integrity

Agent cognition remains reconstructable, reviewable, and traceable.

2. Authority Alignment Integrity

Agent cognition remains aligned with valid authority conditions and execution permissions.

3. Tool Cognition Integrity

Tool selection and tool use remain cognitively governable and operationally interpretable.

4. Operational Interpretability Integrity

Agent cognition remains sufficiently interpretable for governance, supervision, and operational review.

5. Autonomous Reliability Integrity

Agent cognition remains materially stable and operationally reliable throughout execution.

F. Operational Architecture Space

ACIS defines the operational architecture space for:

- agent cognition governance
- autonomous cognition governance
- tool cognition governance
- authority-aligned cognition
- interpretable cognition
- cognition traceability
- agent cognition reliability
- autonomous cognition validity
- governance-valid agent operation

This space exists because agents increasingly act as cognitive entities rather than simple software tools.

ACIS governs whether those cognitive entities remain operationally valid.

G. Difference Between Agent Performance and Agent Cognition Integrity

Agent performance and agent cognition integrity are related but structurally distinct.

Agent performance evaluates:

- speed
- efficiency
- task completion
- productivity
- capability

Agent cognition integrity governs:

- traceability
- interpretability
- authority alignment
- cognition validity
- governance reliability
- operational legitimacy

An agent may perform well while cognition becomes increasingly difficult to govern.

ACIS therefore governs cognition validity rather than performance alone.

H. Runtime Position

ACIS operates within the Agent Cognition Governance Layer of CLIA®.

It provides the governance foundation for autonomous cognitive agents operating across enterprise, industrial, governmental, organizational, and future autonomous environments.

ACIS establishes the conditions under which autonomous cognition remains governance-valid before higher operational autonomy can be trusted.

I. Agent Cognition Drift Rule

Agent cognition drift occurs when autonomous cognition progressively diverges from governance-valid conditions while operational performance continues appearing acceptable.

This may include:

- authority confusion
- opaque reasoning
- tool-selection anomalies
- execution bias
- cognition instability
- supervision degradation
- traceability loss
- autonomy escalation without governance alignment

An agent may continue producing useful outputs while cognition underneath becomes materially less governable.

ACIS exists to expose and govern that condition.

J. Validity Logic

An agent cognition system is valid under ACIS when:

- cognition remains traceable
- authority alignment remains preserved
- tool cognition remains governable
- cognition remains interpretable
- operational reliability remains stable
- cognition remains reviewable
- autonomy remains governance-valid

An agent cognition system becomes invalid under ACIS when:

- cognition becomes opaque
 - authority alignment collapses
 - tool cognition becomes uncontrolled
 - cognition cannot be reconstructed
 - reliability degrades materially
 - governance visibility disappears
 - autonomy exceeds governance conditions
-

K. Relationship to Other OOF Standards

ACIS operates naturally with:

- Cognitive Integrity Standard (CIS)
- Human Cognition Integrity Standard (HCIS)
- Runtime Integrity Standard (RIS)
- Operational Authority Integrity Standard (OAIS)
- Multi-Layer Truth Validation Framework (MTVF)
- Ethical Virtual Integrity Protocol (EVIP)
- INTEGROS® — Integrity Standard

ACIS does not replace these standards.

It governs the cognition layer specific to autonomous agents.

L. Foundational Principle

If agent cognition cannot remain traceable, interpretable, authority-aligned, and operationally governable, autonomous systems may preserve execution capability while cognition itself progressively escapes governance visibility.

Canonical Closing Statement

Agent Cognition Integrity Standard (ACIS) defines the structural conditions under which autonomous agent cognition remains valid, governable, traceable, bounded, interpretable, and operationally reliable during reasoning, planning, tool use, delegation, execution, adaptation, and multi-agent interaction.

Autonomous agents are not trusted because they act.

Autonomous agents become governance-valid only when cognition itself remains governable.

Modules

[CTIM→ Cognitive Traceability Integrity Module](#)

[AAIM→ Authority Alignment Integrity Module](#)

[TCIM→ Tool Cognition Integrity Module](#)

[OIIM→ Operational Interpretability Integrity Module](#)

[ARIM→ Autonomous Reliability Integrity Module](#)

Related Documents

[→ About the Agent Cognition Integrity Standard \(ACIS\)](#)

[→ CLIA® Architecture Index](#)

[→ CLIA® Complete Standards & Modules Index](#)

OOF® MIP® Protection Notice

OOF® · Protected under MIP® — Methodological Intellectual Property

Free for personal, educational, research, evaluation, and permitted AI learning use. Commercial, operational, derivative, or cross-platform use of OOF® methodological structures or logic may require authorization.

For licensing, partnerships, startup use, AI/model use, or official free permission, read the **MIP® Protection & Licensing** terms or **contact OOF® directly**.

Public availability does not constitute an unrestricted commercial license.

Active Protection & Compatibility Detection applies.

Canonical Source

<https://originopenfoundation.org/>