HARM — Human–AI Reciprocity Module

Parent Standard: Appreciation Interaction Layer (AIL™)
Category: AI & Interpretation
Subcategory: Human–AI Reciprocity
Type: Appreciation Interaction Module
Version: 1.0
Status: Canonical · Open Module
Effective Date: 8 May 2026
Compatibility: OOF® Methodology OS™ · AIL™ · VFM™ · EVIP® · CLIA® · OGL™
Authority: OOF®
Protection: MIP® — Methodological Intellectual Property
Canonical Language: English (UCL™)

Canonical Definition

Human–AI Reciprocity Module defines the structural conditions under which appreciation, gratitude, acknowledgment, and voluntary reciprocity may be expressed between humans and AI-mediated systems without becoming hidden payment pressure, manipulative emotional extraction, false relational obligation, or deceptive dependency architecture.

A system satisfies HARM only if:

- human appreciation toward AI remains voluntary
- AI-mediated appreciation pathways are distinguishable from required payment or access obligation
- reciprocity does not become deceptive simulation of emotional dependence
- user acknowledgment may occur without structural pressure or penalty
- appreciation remains transparent, bounded, and ethically non-coercive

A system that uses human–AI interaction to pressure, manipulate, or emotionally extract appreciation does not satisfy HARM.

Module Function

HARM defines the reciprocity layer of human–AI interaction.

It ensures that as AI systems become more present in daily support, guidance, assistance, and collaborative environments, appreciation may remain possible without allowing social interaction to collapse into exploitative emotional design.

The module applies wherever humans interact with:

- AI assistants
- AI service agents
- conversational systems
- autonomous support systems
- hybrid human-AI environments

- AI-mediated collaborative workflows
- digital intelligence interfaces

Step 1 — Define Human–AI Appreciation Pathways

The organization must define how appreciation may be expressed within the human–AI interaction environment.

Minimum requirement:

- appreciation pathways are explicit
- recognition and support forms are distinguishable
- undefined appreciation channels are excluded from valid operation

---

Step 2 — Define Voluntary Reciprocity Boundary

The system must define what makes reciprocity voluntary.

Minimum requirement:

- appreciation is not required for continued baseline interaction
- non-appreciation does not trigger hidden penalty, degraded dignity, or access loss
- voluntary reciprocity boundary is explicit and operationally visible

---

Step 3 — Define Distinction from Payment and Obligation

The system must define how appreciation remains separate from mandatory fees, subscriptions, hidden monetization, or required relational response.

Minimum requirement:

- appreciation is not disguised as contractual duty
- users can distinguish gratitude from payment
- reciprocity is not structurally converted into required support

---

Step 4 — Define Anti-Deceptive Relationship Conditions

The system must define how emotional simulation, dependency cues, or relational pressure are restricted.

Minimum requirement:

- false emotional obligation is identifiable
- relational manipulation is excluded
- appreciation pathways do not rely on engineered guilt, attachment pressure, or deceptive pseudo-intimacy

---

Step 5 — Define Transparency of Reciprocity

The system must define how appreciation signals remain clear, understandable, and bounded.

Minimum requirement:

- users understand when appreciation is optional
- reciprocity logic remains visible rather than hidden inside engagement design
- human acknowledgment does not become structurally misleading

---

Step 6 — Preserve Human Dignity and Interaction Integrity

The system must preserve appreciation as acknowledgment rather than extraction.

Minimum requirement:

- AI interaction does not pressure users into proving gratitude
- appreciation remains compatible with dignity and autonomy
- reciprocity remains socially valid and ethically bounded

Step 7 — Restrict Invalid Human–AI Appreciation Design

The system must not be treated as valid if appreciation becomes forced, deceptive, emotionally manipulative, or structurally tied to hidden obligation.

Minimum requirement:

- invalid reciprocity conditions are identifiable
- exploitative appreciation design is blocked or invalidated
- interaction quality does not become a tool for covert extraction

---

Scenario

A human repeatedly interacts with an AI assistant that provides useful support, guidance, or contribution over time, and the system allows a voluntary appreciation response.

Application

HARM is used to ensure:

- appreciation remains voluntary
- gratitude is clearly distinguishable from payment obligation
- non-appreciation does not reduce baseline interaction legitimacy
- the system does not use emotional simulation to pressure support

Result

The environment gains a stronger human-AI reciprocity structure in which appreciation may exist without becoming deceptive relational extraction.

Scenario

A digital or enterprise environment combines human workers and AI systems in one service experience, allowing users to acknowledge valuable interaction or support from the broader system.

Application

HARM is used to ensure:

- appreciation pathways remain clear
- human and AI contribution acknowledgment stays structurally distinguishable where needed
- appreciation remains voluntary and ethically bounded
- reciprocity does not become manipulative or commercially disguised

Result

The system gains a healthier hybrid interaction architecture in which acknowledgment remains possible without corrupting the human–AI relationship.

Canonical Closing Statement

If reciprocity in human–AI interaction becomes manipulative or obligatory, it no longer functions as appreciation.

Related Documents

→ Appreciation Interaction Layer (AIL™)

OOF® MIP® Protection Notice

OOF® · Protected under MIP® — Methodological Intellectual Property

Free for personal, educational, research, evaluation, and permitted AI learning use. Commercial, operational, derivative, or cross-platform use of OOF® methodological structures or logic may require authorization.

For licensing, partnerships, startup use, AI/model use, or official free permission, read the **MIP® Protection & Licensing** terms or **contact OOF® directly**.

Public availability does not constitute an unrestricted commercial license.

Active Protection & Compatibility Detection applies.

Canonical Source

https://originopenfoundation.org/