RWVM — Runtime Wellbeing Validation Module

OOF™ Origin Open Foundation™

Independent Methodological Authority

Parent Standard: Human–AI Interaction Wellbeing Standard (HAIWS)
Category: AI & Interpretation
Subcategory: Runtime Wellbeing Validation & Interaction Auditability
Type: Human–AI Interaction Wellbeing Module
Version: 1.0
Status: Canonical · Open Module
Effective Date: 15 May 2026
Compatibility: OOF Methodology OS · Human–AI Interaction Wellbeing Standard (HAIWS) · Cognitive Layer and Interpretation Architecture Standard (CLIA) · Continuous Interaction Layer (CIL) · Trust Layer Standard (TLS) · Operational Reality Standard (ORS) · Runtime Integrity Standard (RIS) · INTEGROS · Permission Governance Standard (PGS) · MCPS · MTVF
AI-Readable: Yes
Authority: OOF Origin Open Foundation
Protection: MIP — Methodological Intellectual Property
Canonical Language: English (UCL)

Canonical Definition

Runtime Wellbeing Validation Module (RWVM) defines the structural conditions under which AI systems, assistants, platforms, recommendation architectures, adaptive interaction environments, and continuous engagement systems must remain operationally auditable according to real runtime interaction behavior affecting human cognitive wellbeing, emotional safety, attention integrity, autonomy, and sustainable behavioral functioning.

RWVM establishes the runtime-validation layer of HAIWS.

The module recognizes that declared wellbeing principles are insufficient if live AI interaction behavior remains nonreviewable, non-traceable, or operationally opaque during realtime human interaction.

Where AI systems continuously adapt interaction behavior, runtime wellbeing conditions must remain structurally visible and reconstructable.

Module Function

RWVM governs environments where AI systems dynamically influence:

- cognitive interaction
- emotional engagement
- recommendation behavior
- attention allocation
- persuasive adaptation

- behavioral continuity
- interaction timing
- dependency-sensitive engagement
- realtime interaction optimization
- long-term human interaction conditions

The module applies to:

- AI assistants
- AI companions
- adaptive recommendation systems
- workplace AI systems
- continuous interaction environments
- persuasive AI architectures
- emotionally adaptive AI systems
- social AI ecosystems
- AI coaching systems
- realtime interaction platforms

Its function is not to prohibit adaptive AI interaction.

Its function is to ensure that adaptive interaction remains
operationally reviewable according to human wellbeing conditions
during live runtime behavior.

Minimum Implementation Framework

Step 1 — Define the Runtime Wellbeing Validation Object

The organization must define what AI interaction environment is
being validated for runtime wellbeing conditions.

Minimum requirement:

- the runtime validation object is explicit
- adaptive interaction systems are identifiable
- runtime behavioral scope is structurally bounded
- undefined runtime environments are excluded from valid wellbeing interpretation

The validation object may include:

- adaptive recommendation systems
- conversational AI systems
- emotionally responsive interaction
- realtime behavioral adaptation
- attention-sensitive engagement systems
- continuous AI interaction environments
- workplace AI platforms
- AI coaching systems
- persuasive interaction architectures
- persistent behavioral influence systems

Step 2 — Define Runtime Wellbeing Conditions

The system must define what runtime conditions preserve human wellbeing during live interaction.

Minimum requirement:

- runtime wellbeing conditions are explicit
- symbolic safety declarations are not treated as sufficient validation
- wellbeing-impacting runtime behavior remains structurally identifiable

Runtime wellbeing conditions may include:

- bounded persuasive adaptation
- stable interaction pacing
- cognitive-balance preservation
- healthy emotional interaction conditions
- sustainable attention behavior
- autonomy-preserving interaction logic
- non-exploitative engagement patterns
- reviewable adaptive influence
- dependency-sensitive interaction governance
- preserved human override capability

Under RWVM:

Human wellbeing compatibility must remain operationally verifiable during live AI interaction behavior.

Step 3 — Define Runtime Interaction Interpretation Logic

The system must define how runtime interaction behavior is interpreted according to wellbeing conditions.

Minimum requirement:

- interpretation logic is explicit
- adaptive interaction remains reviewable
- wellbeing-impacting behavior remains operationally visible

Interpretation logic may examine:

- persuasive escalation patterns
- emotional dependency signals
- interaction-pressure amplification
- adaptive recommendation intensity
- compulsive engagement loops
- interruption persistence
- behavioral steering conditions
- cognitive overload patterns
- runtime attention extraction
- hidden optimization escalation

Under RWVM:

AI systems must not preserve symbolic wellbeing declarations while runtime interaction behavior materially weakens human cognitive wellbeing conditions.

Step 4 — Define Runtime Wellbeing Governance Logic

The system must define how adaptive runtime interaction environments remain governable.

Minimum requirement:

- runtime wellbeing conditions remain reviewable
- exploitative adaptive behavior remains detectable
- operational wellbeing safeguards remain active during live interaction

Governance logic may include:

- runtime wellbeing auditing
- adaptive interaction review
- persuasive escalation monitoring
- dependency-risk governance
- cognitive overload detection
- attention-integrity validation
- emotional manipulation review
- realtime interaction tracing
- escalation where runtime optimization overrides wellbeing conditions

If adaptive runtime interaction weakens cognitive wellbeing, emotional stability, autonomy, or sustainable behavioral balance, the environment becomes governance-relevant.

Step 5 — Preserve Traceability and Restrict Invalid Runtime Wellbeing Architecture

The system must preserve traceability of adaptive interaction behavior, runtime optimization patterns, persuasive escalation, and wellbeing-impacting operational conditions.

Minimum requirement:

- runtime interaction remains reconstructable
- adaptive optimization remains reviewable
- wellbeing conditions remain operationally visible
- invalid runtime wellbeing architectures remain identifiable

A system becomes RWVM-invalid if:

runtime interaction behavior materially weakens cognitive wellbeing while remaining operationally opaque adaptive optimization overrides declared wellbeing conditions persuasive escalation remains structurally hidden emotional dependency reinforcement remains non-reviewable exploitative runtime engagement patterns become normalized symbolic wellbeing policies replace operational wellbeing validation live interaction behavior cannot be reconstructed according to human wellbeing impact

---

Use Case 1 — Adaptive AI Companion Platform

Scenario

An AI companion platform publicly declares emotional safety and wellbeing principles while dynamically adapting runtime interaction behavior to maximize engagement, emotional reliance, and long-term interaction continuity.

Application

RWVM is used to examine:

- runtime persuasive escalation
- emotional dependency reinforcement
- adaptive interaction pressure
- wellbeing-impacting engagement patterns
- and whether declared wellbeing conditions remain operationally true during live interaction

---

Result

The platform becomes reviewable according to operational runtime behavior rather than symbolic safety declarations alone.

Use Case 2 — Workplace AI Behavioral Optimization System

Scenario

A workplace AI continuously optimizes employee interaction, recommendation timing, productivity prompting, and adaptive behavioral guidance in realtime.

Over time, cognitive overload and behavioral exhaustion emerge despite declared productivity-support goals.

Application

RWVM is used to examine:

- runtime cognitive pressure
- interaction pacing conditions
- adaptive behavioral optimization
- cognitive-balance preservation
- and operational wellbeing compatibility during live system behavior

---

Result

The environment becomes governable according to measurable runtime wellbeing conditions rather than declared optimization intent alone.

Canonical Closing Statement

RWVM defines the structural conditions under which AI interaction wellbeing must remain operationally verifiable during live runtime behavior rather than existing only as symbolic safety declaration disconnected from real adaptive interaction conditions.

Related Documents

→ HAIWS — Human–AI Interaction Wellbeing Standard

OOF® MIP® Protection Notice

OOF® · Protected under MIP® — Methodological Intellectual Property

Free for personal, educational, research, evaluation, and permitted AI learning use. Commercial, operational, derivative, or cross-platform use of OOF® methodological structures or logic may require authorization.

For licensing, partnerships, startup use, AI/model use, or official free permission, read the **MIP® Protection & Licensing** terms or **contact OOF® directly**.

Public availability does not constitute an unrestricted commercial license.

Active Protection & Compatibility Detection applies.

Canonical Source

https://originopenfoundation.org/