MAACM — Multi-Agent Authority Alignment Module

Parent Standard: Orchestration Governance Layer (OGL™)
Category: AI & Interpretation
Subcategory: Multi-Agent Authority Alignment
Type: Orchestration Governance Module
Version: 1.0
Status: Canonical · Open Module
Effective Date: 8 May 2026
Compatibility: OOF® Methodology OS™ · OGL™ · **CLIA®** · **CML** · **RIS** · **INTEGROS®** · **EVIP®** · **Harness Architectures**
Authority: OOF®
Protection: MIP® — Methodological Intellectual Property
Canonical Language: English (UCL™)

Canonical Definition

Multi-Agent Authority Alignment Module defines the governance conditions under which authority, decision scope, override rights, escalation priority, and execution control remain explicitly aligned across multiple agents, models, tools, runtimes, or orchestration entities.

A system satisfies MAACM only if:

- authority roles are explicitly defined
- decision scope is bounded for each participating entity
- override and escalation rights remain identifiable
- authority conflict is resolved under governed conditions
- distributed execution does not proceed under ambiguous, overlapping, or hidden authority structures

A system that coordinates multiple agents without governed authority alignment does not satisfy MAACM.

Module Function

MAACM defines the authority-governance boundary of orchestration.

It ensures that multi-agent coordination does not create execution without accountable structure, clear hierarchy, or controlled override logic.

The module applies wherever multiple autonomous or semi-autonomous entities participate in distributed decision-making, delegated execution, escalation flow, or shared operational control.

Minimum Implementation Framework (MIF)

Step 1 — Define Authority Roles

The organization must define which entities hold which authority positions.

Minimum requirement:

- authority roles are explicit
- decision-bearing and non-decision-bearing entities are distinguishable
- undefined authority is excluded from valid orchestration logic

---

Step 2 — Define Decision Scope

The system must define what each authority-bearing entity may decide, approve, restrict, or reject.

Minimum requirement:

- scope boundaries are explicit
- authority is not treated as open-ended
- invalid scope expansion is excluded from valid operation

Step 3 — Define Override and Escalation Rights

The system must define who may override, escalate, suspend, or terminate execution.

Minimum requirement:

- override rights are explicit
- escalation hierarchy is identifiable
- hidden override power is excluded from valid orchestration design

---

Step 4 — Define Conflict Resolution Logic

The system must define how conflicting decisions, contradictory instructions, or overlapping authorities are resolved.

Minimum requirement:

- conflict-resolution conditions are explicit
- unresolved authority contradiction does not silently pass into execution
- arbitration or priority logic is defined

Step 5 — Preserve Authority Traceability

The system must preserve visibility of who held authority, who acted, and under which governance condition action was allowed.

Minimum requirement:

- authority path is traceable
- decision origin remains visible
- later review can reconstruct who authorized, overrode, escalated, or restricted execution

---

Step 6 — Restrict Invalid Authority Operation

The system must not be treated as valid if execution proceeds under ambiguous, bypassed, contradictory, or untraceable authority conditions.

Minimum requirement:

- invalid authority conditions are identifiable
- execution under unresolved authority ambiguity is blocked
- coordination convenience does not override authority integrity

---

Use Case 1 — Enterprise Multi-Agent Decision Workflow

Scenario

An enterprise AI environment uses multiple agents for analysis, planning, recommendation, approval support, and execution triggering across one workflow.

Application

MAACM is used to ensure:

- each agent's authority role is explicit
- decision scope remains bounded
- override and escalation rights are defined
- conflicting agent recommendations do not silently become execution authority

---

Result

The organization gains a stronger basis for governing distributed decision flow without allowing multi-agent coordination to blur accountability or decision legitimacy.

Use Case 2 — Dynamic Runtime Coordination with Escalation

Scenario

A runtime orchestration system dynamically shifts tasks between agents while certain conditions allow escalation, override, or execution restriction during live operation.

Application

MAACM is used to ensure:

- authority hierarchy remains explicit during dynamic coordination
- escalation rights are traceable
- override conditions are governed
- execution does not continue under hidden or conflicting authority structures

---

Result

The organization gains a structurally governed authority layer instead of relying on emergent coordination behavior as if it were automatically legitimate.

Canonical Closing Statement

If multiple agents act without aligned authority, orchestration integrity has not been achieved.

Related Documents

→ Orchestration Governance Layer (OGL™)

OOF® MIP® Protection Notice

OOF® · Protected under MIP® — Methodological Intellectual Property

Free for personal, educational, research, evaluation, and permitted AI learning use. Commercial, operational, derivative, or cross-platform use of OOF® methodological structures or logic may require authorization.

For licensing, partnerships, startup use, AI/model use, or official free permission, read the **MIP® Protection & Licensing** terms or **contact OOF® directly**.

Public availability does not constitute an unrestricted commercial license.

Active Protection & Compatibility Detection applies.

Canonical Source

https://originopenfoundation.org/