

ATRIM — Agent Trust & Risk Integrity Module

OOF™ Origin Open Foundation™

Independent Methodological Authority

Modul1 ATRIM — Agent Trust & Risk Integrity Module

OriginID: OOF-OID-GOV-ATRIM-2026-06-02-0001

Architecture Ecosystem: Structured Reality Standards™

Architecture Family: Operational Reality Standards™

Operational Layer: Autonomous Systems Governance Layer

Governed Space: Agent Trust & Risk Integrity

Category: Governance & Enforcement

Subcategory: Agent Trust & Risk Governance

Type: Operational Risk & Trust Integrity Module

Parent Standard: Operational Risk & Trust Integrity Standard (ORTIS)

Version: 1.0

Status: Canonical · Open Module

Origin Date: 2 June 2026

Compatibility: OOF Methodology OS · Operational Risk & Trust Integrity Standard (ORTIS) · Runtime Integrity Standard

(RIS) · Operational Authority Integrity Standard (OAIS) · Operational Decision Integrity Standard (ODIS) · Operational Evidence & Auditability Standard (OEAS) · Operational Escalation Integrity Standard (OESIS) · INTEGROS® — Integrity

Standard · Autonomous Agents · Multi-Agent Systems

AI-Readable: Yes

Authority: OOF

Protection: MIP — Methodological Intellectual Property

Canonical Language: English (UCL)

Canonical Definition

Agent Trust & Risk Integrity Module (ATRIM) defines the structural conditions under which agent trustworthiness, runtime risk exposure, execution eligibility, authority-dependent permissions, and consequence-bearing agent actions remain materially stable, traceable, governable, and operationally valid across autonomous operational environments. ATRIM governs the trust-risk relationship of autonomous agents. The module ensures that agent permissions, execution rights, and operational capabilities remain aligned with current trust conditions and runtime risk states.

A system satisfies ATRIM only if:

- agent trust remains materially assessable
- runtime risk remains governable
- execution eligibility remains contextually valid
- trust degradation remains detectable
- risk escalation remains traceable
- consequence-bearing actions remain accountable
- A system that continues granting execution rights despite material trust degradation or unacceptable risk conditions does not satisfy ATRIM.

Module Operational Role

ATRIM defines the agent trust and risk governance layer of ORTIS. Its role is to preserve governance-valid execution eligibility for autonomous agents operating under changing trust and risk conditions.

Module Operational Space

- ATRIM governs:
 - agent trustworthiness
 - runtime risk exposure
 - execution eligibility
 - trust degradation
 - trust recovery
 - risk escalation
 - permission legitimacy
 - consequence-bearing agent execution
- The module applies wherever autonomous agents perform actions capable of producing operational consequences.

Module Function

- The module applies wherever systems must preserve:
 - trust-valid execution
 - risk-aware permissions
 - accountable agent actions
 - authority-aligned execution rights
 - consequence-bearing operational legitimacy
 - governance-valid runtime behavior
- Its function is to ensure that agent execution rights remain materially aligned with current trust and risk conditions.

Minimum Implementation Framework

1. Define the Agent Trust & Risk Object

The organization must define which agent activities require trust and risk governance. This may include: autonomous task execution tool usage workflow automation delegated authority actions external system interactions consequence-bearing operational decisions multi-agent coordination activities

2. Define Trust & Risk Conditions

The system must define the conditions under which agent trust and risk remain operationally valid. This includes: trust thresholds risk thresholds permission eligibility conditions authority requirements sensitivity conditions consequence-bearing execution limits

3. Define Trust Degradation & Risk Escalation Detection Logic

The system must define how trust degradation and risk escalation are identified. This may include: abnormal behavior detection permission misuse detection operational anomaly detection authority violations elevated risk conditions execution legitimacy degradation

4. Define Operational Response or Governance Logic

The system must define governance logic for materially unstable trust-risk conditions. Governance response may include: permission restriction execution limitation trust re-evaluation

- risk escalation review
- operational intervention
- execution suspension
- operational invalidation where required

5. Preserve Traceability & Restrict Invalid Conditions

- The system must preserve reconstructable traceability of:
 - trust state changes
 - risk state changes
 - permission decisions
 - execution eligibility decisions
 - escalation events
 - consequence-bearing agent actions
- A system must not remain trust-valid if agent execution continues despite materially degraded trust conditions or
- unacceptable risk exposure.

Use Case 1 — Enterprise Autonomous Agent

Platform

Scenario

An enterprise deploys autonomous agents capable of accessing internal systems, coordinating workflows, and executing operational tasks across multiple business environments.

Application

ATRIM continuously governs trust levels, runtime risk exposure, and execution eligibility based on current operational behavior and authority conditions.

Result

The organization gains stronger agent autonomy while preserving governance-valid execution permissions and risk visibility.

Use Case 2 — Multi-Agent Financial Operations

Environment

Scenario

A distributed multi-agent environment performs autonomous financial analysis, transaction preparation, and operational coordination across sensitive financial systems.

Application

ATRIM governs whether agents remain eligible to perform actions based on trust conditions, authority legitimacy, and runtime risk exposure.

Result

The environment gains stronger operational accountability while reducing unauthorized execution and uncontrolled risk escalation.

Canonical Closing Statement

Agent Trust & Risk Integrity Module (ATRIM) defines the structural conditions under which agent trustworthiness, runtime risk exposure, execution eligibility, authority-dependent permissions, and consequence-bearing agent actions remain materially stable, traceable, governable, and operationally valid across autonomous operational environments. Autonomous agents are not valid merely because they possess permission to act. Agent execution becomes valid only when trust, risk, authority, and consequence-bearing operational conditions remain materially aligned throughout runtime operation.

OOF® MIP® Protection Notice

OOF® · Protected under MIP® — Methodological Intellectual Property

Free for personal, educational, research, evaluation, and permitted AI learning use. Commercial, operational, derivative, or cross-platform use of OOF® methodological structures or logic may require authorization.

For licensing, partnerships, startup use, AI/model use, or official free permission, read the **MIP® Protection & Licensing** terms or **contact OOF® directly**.

Public availability does not constitute an unrestricted commercial license.

Active Protection & Compatibility Detection applies.

Canonical Source

<https://originopenfoundation.org/>