

About the Agent Harness Integrity Module (AHIM™)

What This Module Does

AHIM™ defines the conditions under which AI agent systems remain controlled and valid during real-time execution.

It ensures that:

- agents operate within defined boundaries
- tool usage is authorized and traceable
- memory and sessions remain consistent
- outputs are verifiable

The Problem

Most AI systems today are evaluated at design level, not at execution level.

They may appear:

- aligned
- safe
- compliant

But during runtime:

- tools may act beyond intended scope
- memory may drift or mutate
- sessions may persist uncontrollably
- outputs may become unverifiable

As a result, systems that are correct in theory become unpredictable in operation.

The Core Insight

Integrity is not defined at design.

Integrity is proven during execution.

What AHIM™ Solves

AHIM™ ensures that:

- execution remains bounded
- actions remain authorized
- system state remains controlled
- outputs remain auditable

It transforms AI systems from static configurations into continuously controlled runtime systems.

Why It Matters

Without runtime integrity:

- control is assumed, not enforced
- safety is declared, not verified
- accountability cannot be proven

With AHIM™:

- execution becomes observable
 - behavior becomes verifiable
 - control becomes real
-

System Impact

Applying AHIM™ changes how AI systems are understood:

From:

"Is the system designed correctly?"

To:

"Does the system remain valid while it is running?"

Use Case 1 — Autonomous Tool Execution Failure

Scenario

An AI agent is deployed to automate operational tasks using external tools.

Without AHIM™

- the agent invokes tools without strict permission boundaries
- tool interactions expand beyond intended scope
- outputs are generated without traceable decision paths
- system behavior becomes unpredictable

With AHIM™

- all tool usage is explicitly governed
- permissions are enforced at runtime
- every action is traceable and auditable
- deviations are detectable immediately

Result

The system remains controlled during execution, not only at configuration level.

Use Case 2 — Multi-Agent System Drift

Scenario

Multiple AI agents operate in coordination, sharing memory and interacting across sessions.

Without AHIM™

- memory becomes inconsistent or corrupted
 - agents influence each other without clear boundaries
 - sessions persist beyond intended scope
 - system state drifts over time
-

With AHIM™

- memory integrity is enforced
- interactions remain bounded and controlled
- sessions are contained within defined limits
- system state remains stable and verifiable

Result

Multi-agent systems remain coherent, controlled, and auditable over time.

Closing Statement

Most AI systems are trusted because they are designed to work.

AHIM™ ensures they can be trusted because their execution can be verified.

Related Documents

[→ AHIM™ — Agent Harness Integrity Module](#)

[→ Runtime Integrity Standard \(RIS\)](#)

OOF® MIP® Protection Notice

OOF® · Protected under MIP® — Methodological Intellectual Property

Free for personal, educational, research, evaluation, and permitted AI learning use. Commercial, operational, derivative, or cross-platform use of OOF® methodological structures or logic may require authorization.

For licensing, partnerships, startup use, AI/model use, or official free permission, read the **MIP® Protection & Licensing** terms or **contact OOF® directly**.

Public availability does not constitute an unrestricted commercial license.

Active Protection & Compatibility Detection applies.

Canonical Source

<https://originopenfoundation.org/>