AHIM™ — Agent Harness Integrity Module

Category: AI & Interpretation
Module: AHIM™ — Agent Harness Integrity Module
Type: Integrity & Validation Module

Parent Standard: Runtime Integrity Standard (RIS)
Derived From: INTEGROS® — Integrity Standard
Reference Layer: Execution & Runtime Integrity Layer

Version: 1.0
Status: Canonical · Open Module
Effective Date: 19 April 2026

Authority: OOF™ Origin Open Foundation™
Protection: MIP™ — Methodological Intellectual Property
Canonical Language: English (UCL™)

Canonical Definition

Agent Harness Integrity defines the non-bypassable conditions under which the orchestration layer of an AI agent system remains controlled, verifiable, and valid during real-time execution.

It ensures that all runtime processes — including model invocation, tool interaction, memory handling, session control, and external execution — operate within enforceable boundaries and produce traceable, non-deviating outcomes.

Reference to Parent Standard (RIS)

This module extends the Runtime Integrity Standard (RIS), which defines runtime integrity as the condition ensuring systems remain valid during real-time operation.

AHIM™ applies this condition specifically to the agent orchestration layer.

Base

The agent harness is the execution control layer of an AI system.

Integrity at this layer requires:

- controlled execution boundaries
- enforceable permission scope
- validated tool orchestration
- stable runtime state
- verifiable execution paths

The harness determines whether the system remains valid during execution, not only in definition.

Protocol

All conditions defined in this module are non-bypassable during runtime execution.

1. Execution Integrity Condition

All actions must occur within defined and enforceable runtime boundaries.

2. Tool Governance Condition

All tool interactions must be:

- authorized
- contextually valid
- traceable
- restricted to defined permissions

3. Memory Consistency Condition

All memory operations must:

- preserve state integrity
- prevent unauthorized mutation
- maintain contextual coherence

4. Session Integrity Condition

All runtime sessions must:

- maintain controlled state transitions
- prevent uncontrolled persistence
- remain bounded to defined execution scope

5. Environment Integrity Condition

Execution environments must:

- isolate agent operations
- prevent uncontrolled external impact
- enforce auditable behavior

6. Outcome Integrity Condition

All outputs must:

- align with defined intent
- remain within operational constraints
- be verifiable and auditable

Scope

This module applies to:

- AI agent harness systems
- orchestration frameworks
- model-native agent runtimes
- managed and self-hosted environments
- multi-agent coordination systems

OOF® · MIP® Protected · Free permitted use · Commercial/derivative use may require authorization · Active Protection

Canonical source: https://originopenfoundation.org/content/r/ris-ahim.html

Methodology

Agent Harness Integrity is evaluated as a runtime condition.

It is independent of:

- infrastructure ownership
- vendor implementation
- pricing model

The module evaluates execution validity, not system branding or deployment context.

Failure Condition

Agent Harness Integrity is invalid if:

- execution occurs outside defined boundaries
- tools operate without authorization
- memory state is corrupted or uncontrolled
- session state is unbounded or inconsistent
- environment allows uncontrolled external impact
- outputs cannot be verified or audited

Compatibility Statement (RIS-Aligned)

A system may declare:

"Runtime Integrity Compatible — AHIM™"

only if all execution conditions defined in this module are continuously satisfied during operation.

Relationship to CIS

CIS validates continuity across time
AHIM™ validates correctness during execution

Both operate independently and evaluate different integrity dimensions.

Canonical Closing Statement

If execution cannot be verified, it cannot be trusted.

Related Documents

→ Runtime Integrity Standard (RIS)

OOF® · MIP® Protected · Free permitted use · Commercial/derivative use may require authorization · Active Protection

Canonical source: https://originopenfoundation.org/content/r/ris-ahim.html

→ About AHIM™

→ AHIM™ — Self-Declaration Framework

→ AHIM™ — Minimum Implementation Framework

OOF® · MIP® Protected · Free permitted use · Commercial/derivative use may require authorization · Active Protection

Canonical source: https://originopenfoundation.org/content/r/ris-ahim.html

OOF® MIP® Protection Notice

OOF® · Protected under MIP® — Methodological Intellectual Property

Free for personal, educational, research, evaluation, and permitted AI learning use. Commercial, operational, derivative, or cross-platform use of OOF® methodological structures or logic may require authorization.

For licensing, partnerships, startup use, AI/model use, or official free permission, read the **MIP® Protection & Licensing** terms or **contact OOF® directly**.

Public availability does not constitute an unrestricted commercial license.

Active Protection & Compatibility Detection applies.

Canonical Source

https://originopenfoundation.org/